

Regime-Dependent Predictive Structure Between Equity Factors: Evidence from Granger Causality

Anonymous Author(s)

Abstract

We evaluate regime-dependent predictive structure between equity factors using 35 years of Fama-French daily data (1990–2024, 8,817 trading days) and a Student- t Hidden Markov Model with 50 multi-start initializations. To eliminate regime-identification circularity, the *primary* analysis uses a frozen out-of-sample design: the HMM is trained on 1990–2012 ($n = 5,797$ days) and applied without refitting to classify 2013–2024 ($n = 3,020$ days), after which Granger tests are run on the held-out labels. In this frozen OOS evaluation, HML→SMB shows suggestive evidence in the Elevated regime ($p = 0.014$; HAC-robust $p = 0.041$; $\Delta R^2 = 0.73\%$; $n = 836$; not Bonferroni-significant across 30 pairs), with Normal and Crisis null ($p = 0.149$ and $p = 0.281$). A full 50-seed sensitivity analysis finds that 2 of 3 local optima (covering 41 of 50 seeds; F - p and HAC- $p < 0.05$) show Elevated significance; the 9 null seeds all share the lowest-LL local optimum, so effective robustness is across distinct optima rather than independent initializations. Full-sample in-sample analysis corroborates the directional finding with higher power: Normal-regime HML→SMB $p = 8.75 \times 10^{-9}$ (HAC $p = 9.45 \times 10^{-8}$); the shift from Normal (in-sample) to Elevated (OOS) reflects regime redistribution under the frozen classifier: Crisis shrinks from 64.8% (train) to 45.3% (test) while Elevated grows from 13.7% to 30.7%. Event validation shows 2/6 significant, 4/6 directional. Exploratory FF25 mechanism analysis ($\rho_s = 0.35$, $p = 0.046$) is consistent with portfolio overlap. The trading strategy is not profitable (Sharpe = -0.07). A hybrid VaR detector yields 5.60% violations with Christoffersen CC [16] $p = 0.336$. We treat Granger findings as predictive precedence, not structural causality.

CCS Concepts

• **Mathematics of computing** → Time series analysis; • **Computing methodologies** → Causal reasoning and diagnostics.

Keywords

Factor Investing, Granger Causality, Regime Switching, Hidden Markov Models, Model Complexity Characterization, Out-of-Sample Validation, Risk Management

ACM Reference Format:

Anonymous Author(s). 2026. Regime-Dependent Predictive Structure Between Equity Factors: Evidence from Granger Causality. In *Proceedings of* . ACM, New York, NY, USA, 14 pages. <https://doi.org/10.1145/nnnnnnn>. nnnnnnn

1 Introduction

The August 2007 quantitative meltdown, in which systematic equity strategies lost approximately 30% in three days [39], revealed a critical blind spot in factor-based risk management. When multiple quantitative funds held similar factor exposures, forced liquidation

by one fund created price pressure that cascaded to all others. Standard correlation-based risk models failed to anticipate this cascade because they measure *co-movement* but not *temporal precedence*.

The Missing Piece: Which Factor Moves First?

Existing research establishes three stylized facts: (1) Factor correlations increase during market stress [3]; (2) Returns exhibit regime-switching behavior [34]; (3) Factor crowding amplifies drawdowns during liquidation [55].

However, a critical question remains underexplored: **Does the predictive architecture between factors change qualitatively across market regimes?**

Correlation is symmetric—it cannot distinguish whether Value movements precede Size movements or vice versa. Granger causality tests for temporal precedence: whether past values of one series improve prediction of another, conditional on the target’s own history. If such predictive relationships vary by regime, we could:

- Monitor the leading factor to anticipate movements in the lagging factor
- Adjust hedges based on the current regime’s predictive structure
- Detect regime transitions by observing when predictive links emerge or disappear

Important caveat. Granger causality establishes *predictive precedence*, not *structural causality*. A statistically significant Granger relationship could arise from: (1) direct causal influence, (2) a common driver affecting both series with different lags, or (3) omitted variables. We interpret our findings as documenting regime-dependent predictive structure, with economic mechanisms offered as hypotheses rather than verified causal channels. This distinction between predictive precedence and structural causality is foundational in the causal inference literature [48] and has been formalized in the recent factor investing literature [43]; we adopt the Granger framework for its tractability on high-dimensional daily factor series while acknowledging that López de Prado advocates structural causal models (SCMs/DAGs) as the more rigorous alternative for factor research.

1.1 Our Findings

Using Granger causality analysis within regime-dependent subsamples identified by a Student- t Hidden Markov Model, we establish:

Main Finding: regime-dependent predictability with suggestive OOS corroboration. Using a frozen OOS design (HMM trained on 1990–2012, Granger tested on 2013–2024), HML→SMB shows suggestive evidence in stress-adjacent regimes with no circularity (not Bonferroni-significant across 30 pairs; economically motivated pair):

- *Elevated* (OOS, $n = 836$): $p = 0.014$, HAC $p = 0.041$ (Newey–West, bandwidth 1), $\Delta R^2 = 0.73\%$ (lag 1, pre-fixed from in-sample BIC).
- *Normal* (OOS): $p = 0.149$ (null).

- *Crisis (OOS)*: $p = 0.281$ (null).

A full 50-seed sensitivity analysis finds that 2 of 3 local optima (covering 41 of 50 seeds; $F\text{-}p < 0.05$) show Elevated significance; the 9 null seeds all share the lowest-LL local optimum. Full-sample in-sample analysis provides higher-power corroboration: Normal-regime $p = 8.75 \times 10^{-9}$ (HAC $p = 9.45 \times 10^{-8}$), surviving Bonferroni correction across 30 directed pairs. A single-stage Markov-Switching VAR further supports HML lags ($p < 10^{-12}$), and event tests are mixed (2/6 significant, 4/6 in the predicted direction) (Section 5.2).

Predominantly linear. A four-model diagnostic protocol (Linear, RF, MLP, LSTM) finds that no nonlinear method improves HML→SMB prediction beyond the linear baseline. However, transfer entropy reveals an additional nonlinear reverse channel (SMB→HML, $z = 5.37$, $p < 0.005$ empirical) in Normal, invisible to standard Granger tests (Section 5.6).

Mechanism evidence (sensitivity fit, seed 42). Granger significance concentrates in small-cap portfolios within the FF 25 grid ($\rho_s = 0.35$, permutation $p = 0.046$), consistent with portfolio overlap; this analysis uses an alternative HMM fit optimized for crisis detection rather than the primary OOS fit (see Section 4.5).

Practical Implication. The effect is small ($\Delta R^2 \approx 1\%$) and does not translate to profitable trading, but a hybrid VaR model combining HMM regime detection with a realized-volatility fallback achieves 5.60% violation rate ($p_{CC} = 0.336$)—the best among all specifications (Section 5.5).

1.2 Contributions

- (1) **Empirical (circularity-free):** Suggestive OOS evidence for HML→SMB in the Elevated regime (frozen design, 2013–2024, $n = 836$): $p = 0.014$, HAC $p = 0.041$; not Bonferroni-significant across 30 pairs; Normal and Crisis null. Full-sample in-sample analysis corroborates with higher power ($p = 8.75 \times 10^{-9}$, Normal regime, Bonferroni-significant) (Section 5.2, 4.3)
- (2) **Methodological:** Student- t HMM with multi-start selection (50 initializations) for reproducible regime detection (Sections 3.2, 4.2)
- (3) **Complexity Characterization:** A four-model diagnostic protocol (Linear, RF, MLP, LSTM) maps the complexity boundary of the predictive link, establishing the forward HML→SMB link as predominantly linear while identifying a nonlinear reverse-direction channel via transfer entropy (SMB→HML, $z = 5.37$, $p < 0.005$ empirical), demonstrating the protocol’s diagnostic value (Section 5.6)
- (4) **Exploratory Mechanism Evidence:** Portfolio overlap analysis of the FF 25 Size×BM grid provides exploratory evidence consistent with the hypothesis that the HML→SMB link is mediated by small-value portfolios loading on both factors (sensitivity fit, seed 42; spatial gradient test: $\rho_s = 0.35$, permutation $p = 0.046$; asymptotic $p = 0.086$; Section 4.5)
- (5) **Honest Assessment:** Transparent evaluation showing the finding supports risk monitoring rather than alpha generation, with clear acknowledgment of what Granger causality can and cannot establish (Sections 5.4, 5.7)

2 Related Work

Factor models and cross-sectional returns. Fama and French [24, 25] establish that Size, Value, Profitability, and Investment factors explain cross-sectional return variation. Asness et al. [5] document value and momentum effects globally. We study dynamic *interactions* between these factors rather than their individual risk premia.

Factor time-series predictability. Jegadeesh and Titman [38] establish cross-sectional momentum (past winners outperform past losers), with Carhart [11] formalizing MOM as a systematic factor. Ehsani and Linnainmaa [22] show that factor returns exhibit time-series momentum: past factor returns predict future factor returns, with the momentum factor loading on lagged returns across all other factors. Gupta and Kelly [32] extend this finding across asset classes. Haddad, Kozak and Santosh [33] show that individual factor returns are predictable from conditioning variables representing the stochastic discount factor (out-of-sample $R^2 \approx 4\%$)—a within-factor SDF-based timing result. Flögel, Schlag and Zunft [26] exploit cross-factor first-order autocorrelations to construct momentum-managed factor portfolios. Barroso and Santa-Clara [6] and Daniel and Moskowitz [19] document that momentum profits are regime-dependent: high in tranquil periods, negative during momentum-crash episodes. **Gap:** These papers establish within-factor time-series momentum, regime-dependent momentum *returns*, SDF-based factor timing, and cross-factor autocorrelation-based portfolio management (Flögel et al.), but do not examine whether the *Granger predictive structure between distinct factors*—which factor leads the other, and under which regime—changes qualitatively across market states. Our contribution is distinct from factor momentum by econometric construction: the Granger test conditions on SMB’s own lags, ruling out within-factor SMB autocorrelation as an explanation; the HML coefficient tests *incremental* cross-factor content beyond SMB’s history, and our regime conditioning reveals Normal-regime significance invisible to the unconditional autocorrelation approach of Flögel et al. This design reveals regime-dependence invisible to unconditional tests: significant at $p = 8.75 \times 10^{-9}$ in Normal (Bonferroni-significant) but null in Crisis ($p = 0.695$; the Crisis regime contains only 1,071 days, limiting power relative to Normal’s 4,723 days). Additionally, Chordia and Shivakumar [15] document that momentum profits vary systematically with the business cycle (larger in expansions, attenuated or reversed in downturns when macro-risk factors dominate)—the closest single-paper precedent for our regime-conditional predictability finding; we extend this to the cross-factor Granger structure and apply a circularity-free OOS design absent from prior work (Chordia and Shivakumar use observed macroeconomic conditioning, not a latent-state frozen classifier).

Factor crowding and systemic risk. Shleifer and Vishny [51] establish theoretically that capital constraints force arbitrageurs to liquidate positions simultaneously, creating return spillovers across assets. Anton and Polk [4] show that stocks with common mutual fund ownership exhibit correlated returns. Lou and Polk [44] measure arbitrage activity through return comovement. Marks and Shang [45] provide direct empirical evidence that HML/SMB crowding leads to liquidity exhaustion and factor-level crash risk. Chincarini, Lazo-Paz and Moneta [13] show that anomaly returns concentrate among the most crowded stocks and that crowding

positively predicts crash risk, corroborating the HML→SMB contagion channel proposed here. Stein [55] formalizes how crowded trades amplify drawdowns; Brunnermeier and Pedersen [9] establish the theoretical mechanism linking margin constraints to funding-liquidity spirals; Brunnermeier [8] documents the 2007–2008 episode. Coval and Stafford [18] document that forced mutual fund outflows produce asset fire sales with predictable return patterns in affected stocks, providing direct empirical evidence for the liquidation-driven contagion channel. Greenwood and Thesmar [30] show fragility arises from overlapping ownership, providing micro-foundations for cross-factor contagion. **Gap:** Existing work measures crowding *intensity* within individual factors but does not examine *regime-conditional predictive spillover* between factors.

Regime-switching models in finance. Hamilton [34] introduced Markov-switching models for business cycles. Ang and Bekaert [2] apply regime-switching to international asset allocation, establishing standard practice for multi-regime equity analysis. Krolzig [40] extends this to Markov-Switching Vector Autoregressions (MS-VARs), which jointly model regime-dependent dynamics—the single-stage analog of our two-stage approach (we compare against MS-VAR in Section 5.3). Guidolin and Timmermann [31] apply regime models to asset allocation. Chincoli and Guidolin [14] show that Markov-switching VARs modestly improve out-of-sample factor-return predictability for size, value, and momentum, though without examining Granger causal direction between factors. Bulla [10] introduces Student-*t* HMMs for financial returns, motivated by excess kurtosis documented in [17]; Oelschläger, Adam and Michels [47] provide a recent open-source implementation (fHMM) validating this model class on financial data. Shu, Yu and Mulvey [54] compare statistical jump models against HMMs for equity regime detection, finding jump models reduce downside risk in specific settings; Shu and Mulvey [53] apply regime-switching signals to dynamic allocation across six equity factors (value, size, momentum, quality, low-vol, growth) — selecting which factor to hold per regime — whereas the present paper examines Granger predictive *precedence between* factors within shared multi-factor regimes. Our two-stage Granger design motivates the HMM choice for its interpretable latent-state structure. Wang, Lin and Mikhelson [57] combine HMMs with Fama-French factors to switch factor strategies across regimes. **Gap:** Existing regime-switching research focuses on regime-dependent *distributions* or regime-conditional *strategy allocation* (Wang et al. 2020), but does not examine whether Granger predictive *structure between* factors changes across regimes—which is our focus.

Granger causality and connectedness. Lo and MacKinlay [42] establish cross-autocorrelation between large-cap and small-cap portfolios as a foundational lead-lag result; our work extends this cross-portfolio predictability to the factor level and conditions it on latent market regimes. Granger [29] introduced the predictive-precedence test we employ. Psaradakis, Ravn and Sola [49] pioneered regime-switching Granger causality by modeling time-varying causal structure via a Markov chain in a macroeconomic context. Droumaguet et al. [21] extend this to full MS-VAR with Bayesian inference, jointly estimating regime-dependent predictive densities—the single-stage analog of our two-stage approach; we

compare against this benchmark in Section 5.3. We adopt a two-stage approach (HMM then regime-conditional Granger) that allows the regime model and causality test to be estimated separately.¹ Billio et al. [7] construct Granger causality networks among financial institutions to measure systemic connectedness. Diebold and Yilmaz [20] develop directional spillover measures based on forecast error variance decompositions. **Gap:** No prior work examines *regime-conditional* Granger causality between Fama-French equity factors (as opposed to institution-level unconditional networks or regime-conditional strategy allocation).

Neural Granger causality and complexity characterization. Tank et al. [56] extend Granger causality to nonlinear settings using neural networks (MLP, LSTM), enabling detection of nonlinear predictive relationships that linear *F*-tests miss. Howard, Lohre and Mudde [37] apply causal discovery algorithms (DAG-based) to factor investing, constructing unconditional causal networks among factors that are static across time. Their primary finding identifies value and momentum factors as key hubs in the factor causal graph (including value-momentum and value-size linkages)—a result that our HML→SMB (value leads size) finding corroborates for the value-size pair specifically, and extends: we show this predictive precedence is *regime-dependent*, active in Normal conditions but null during Crisis ($p = 0.695$), a distinction their unconditional design cannot detect. We apply the neural Granger framework alongside Random Forest and transfer entropy [28, 50] not to discover new nonlinear links, but to *characterize the complexity boundary* of a known linear relationship—determining whether nonlinear models capture additional predictive structure beyond the linear baseline.

3 Methodology

3.1 Data

We use daily returns for the Fama-French six factors from Kenneth French’s data library [27]: Market excess return (MKT-RF), Size (SMB), Value (HML), Profitability (RMW), Investment (CMA), and Momentum (MOM) [11, 38].

Sample: January 2, 1990 – December 31, 2024 (8,817 trading days). Daily factor returns are stationary by construction; ADF tests confirm unit-root rejection ($p < 0.001$) for all six series.

3.2 Student-*t* Hidden Markov Model

Let $z_t \in \{1, \dots, K\}$ denote the latent regime at time t . We model:

Transition model: $P(z_t = k \mid z_{t-1} = j) = A_{jk}$

Emission model (multivariate Student-*t*):

$$p(\mathbf{x}_t \mid z_t = k) \propto \left(1 + \frac{\delta_k(\mathbf{x}_t)}{v_k}\right)^{-\frac{v_k+d}{2}} \quad (1)$$

where $\delta_k(\mathbf{x}_t) = (\mathbf{x}_t - \boldsymbol{\mu}_k)^\top \Sigma_k^{-1} (\mathbf{x}_t - \boldsymbol{\mu}_k)$.

¹For recent developments in MS-VAR forecasting, see Cavicchioli [12]. The Bayesian MS-VAR of Droumaguet et al. (2017) has a finite-sample advantage for small-regime subsets: posterior inference is exact regardless of subsample size. Our asymptotic *F*-tests with HAC correction partially address finite-sample concerns; results are consistent across HAC bandwidths (sensitivity: bandwidth-1 gives Elevated $p = 0.041$, reported in Table 7 caption), supporting the reliability of asymptotic inference given regime sample sizes ($n \geq 800$ for Primary fit in-sample tests).

Why Student- t ? Financial returns exhibit excess kurtosis. Gaussian HMMs calibrate thresholds to extreme historical observations, missing moderate crises. Student- t distributions accommodate heavy tails, enabling detection of crises with volatility below historical extremes. Estimated degrees of freedom for the primary fit (seed 28): $\hat{\nu}_{\text{Normal}} = 6.2$, $\hat{\nu}_{\text{Elevated}} = 3.9$, $\hat{\nu}_{\text{Crisis}} = 5.5$ —all consistent with heavy tails well below the Gaussian limit. The Elevated regime ($\hat{\nu} = 3.9 < 4$) has technically infinite kurtosis, reflecting extreme tail behavior in stress-adjacent periods; the HAC-robust Granger tests remain valid since they require no finite-moments assumption beyond what is needed for asymptotic normality of OLS coefficients.

Model selection. The number of regimes ($K = 3$) was selected by the Bayesian Information Criterion, which favored three regimes over two ($\Delta\text{BIC} = \text{BIC}(K=2) - \text{BIC}(K=3) = 847$) or four ($\Delta\text{BIC} = \text{BIC}(K=4) - \text{BIC}(K=3) = 312$). Transition probabilities are estimated (not constrained) via the Expectation-Maximization algorithm.²

3.3 Sample Overlap Consideration

A methodological concern is that regime labels are identified using the full sample, then Granger tests are conducted within those regimes. This means regime assignments are not truly out-of-sample with respect to the returns being tested.

We address this concern in two ways. First, the HMM uses *distributional* properties (mean, covariance, tail behavior) to assign regimes, while Granger tests examine *temporal dynamics* (lagged predictive relationships). These are distinct statistical features, limiting circularity. Second, our event-based validation (Section 5.1) tests whether the in-sample pattern generalizes to specific historical episodes, providing partial out-of-sample assessment.

We adopt the fully rigorous approach as the *primary* validation, following the OOS evaluation standard of Welch and Goyal [58]: the HMM is fitted on 1990–2012 and all parameters are frozen before classifying 2013–2024; Granger tests are then run on the held-out labels using lag = 1 pre-fixed from in-sample BIC (Section 5.2). This frozen OOS design eliminates circularity and yields significant HML→SMB in the Elevated regime of the test period. Additionally, BIC model selection ($K = 3$ vs. $K = 2$ vs. $K = 4$) is verified on training data alone: $K = 3$ achieves BIC = 44,652 vs. $K = 2$ ’s 46,332 ($\Delta\text{BIC} = 1,680$), confirming $K = 3$ is BIC-optimal without accessing OOS data. (Note: the $\Delta\text{BIC} = 847$ reported in Section 3.2 is from full-sample model selection on 1990–2024; the $\Delta\text{BIC} = 1,680$ here is from training-only selection on 1990–2012. Both strongly favor $K = 3$.)

²We fit Student- t HMMs with 50 EM initializations. *Primary selection (leakage-safe):* the fit with highest log-likelihood across all 50 starts, regardless of event detection. Three seeds (28, 20, 6) reach the same full-sample log-likelihood; seed 28 is selected as primary by convention (first in sorted order). Note: when re-trained on 1990–2012 only (frozen OOS), these seeds converge to distinct local optima with different train-period log-likelihoods. This fit assigns 0% of 2008, 2011, and 2020 stress windows to the Crisis regime (Table 2). *Sensitivity:* among the 16/50 fits (32%) that assign $\geq 50\%$ of 2008 days to Crisis, we separately report the highest-LL fit; the LL cost is $< 0.3\%$. All tables use the primary fit unless explicitly noted. The FF 25 portfolio overlap analysis (Section 4.5) and VaR monitoring (Section 5.5) use the sensitivity fit (seed 42), which assigns 47.8% of COVID-2020 days to Crisis and yields qualitatively identical full-sample Granger results to the primary fit (see FF25 footnote and VaR section for frozen OOS Granger disclosure); this is disclosed at each analysis. Table 2 compares against a Gaussian HMM baseline (also multi-start, 10 seeds).

3.4 Per-Regime Granger Causality

Granger causality [29, 52] tests whether past values of one series improve prediction of another, conditional on the target’s own history—a *predictive precedence* criterion, not structural causality. For each regime k , we extract observations: $\mathcal{T}_k = \{t : \hat{z}_t = k\}$. For each ordered pair of factors (i, j) , we test:³ $H_0 : r_{j,t} \perp \{r_{i,t-\ell}\}_{\ell=1}^L \mid \{r_{j,t-\ell}\}_{\ell=1}^L$

Significance threshold: $\alpha = 0.01$ with Bonferroni correction across 30 directed factor pairs (effective $\alpha = 0.00033$). HAC standard errors use Newey–West with bandwidth equal to the Granger lag (bandwidth 1) as the primary specification, applied consistently across all Granger tests. As a robustness check, Andrews’s [1] AR(1) plug-in data-driven bandwidth (≈ 9 – 10 for $n \approx 800$) yields directionally consistent results at conventional significance levels for the Elevated OOS regime.

Regime boundary handling. When computing lagged variables for Granger tests, we include only observations where all lags fall within the same regime. Specifically, for an observation at time t in regime k with maximum lag L , we require $\hat{z}_{t-\ell} = k$ for all $\ell \in \{1, \dots, L\}$. The exclusion varies by analysis: Table 3 uses BIC-selected lags (lag 1 for all rows); Table 4 uses 15 lags; Tables 14 and 15 use 9 lags; Table 10 uses model-specific warm-up requirements. For the lag-1 Granger analysis, exclusion also removes observations lacking a consecutive-trading-day predecessor (e.g., Mondays and post-holiday days), which accounts for the majority of the 22% rate. The higher-lag neural and transfer entropy analyses use trading-day-indexed lags and apply regime-consistency checks only, yielding lower per-table exclusion rates despite longer lag windows (e.g., Normal: $n = 4,496$ at lag 9 vs. $n \approx 3,680$ under lag-1 calendar exclusion).

Quantified impact. With lag-1 Granger tests, 22.2% of observations (1,956 of 8,817) are excluded at regime boundaries—distributed as Normal: 22.1% (1,043/4,723), Elevated: 22.4% (676/3,023), Crisis: 22.1% (237/1,071). The gap distribution is highly concentrated at short durations: for all three regimes, the median gap is 3 calendar days and the 95th percentile is 4 days (weekend and holiday breaks). Only 1.2–3.3% of inter-regime gaps exceed 20 trading days (≈ 1 month), indicating that the vast majority of “non-consecutive” observations are separated by at most a weekend—not by long regime interruptions that would undermine the economic interpretation of lag-1 prediction. The maximum gaps (1,568 days Normal, 2,423 days Elevated, 4,957 days Crisis) reflect rare long intervals between *separate* regime episodes, not typical gaps. Boundary exclusion rates are uniform across regimes ($\approx 22\%$), indicating no systematic bias against any particular regime.

Limitation: This approach may introduce selection bias by systematically excluding days immediately following regime transitions—precisely when predictive structure may be most dynamic. We view this as a conservative choice that ensures clean within-regime estimation, but acknowledge it may underestimate transition effects.

³We test pairwise (bivariate) Granger relationships. A high-dimensional VAR with all 6 factors and 9 lags would require $6 \times 6 \times 9 = 324$ parameters per regime; with the smallest regime at $n \approx 1,000$ observations this is severely under-identified. Post-double-selection methods [35] could address this in future work.

Table 1: Regime Summary Statistics (1990–2024, full-sample primary fit). Transition probabilities are estimated via Expectation-Maximization. The frozen OOS analysis (Section 5.2) uses a separate fit trained on 1990–2012 only, yielding different regime counts. Elevated and Crisis regimes have similar mean norms (2.03 vs. 2.09) but are distinguished by tail shape ($\nu = 3.9$ vs. 5.5) and covariance structure.

Regime	Days	Prop.	Mean $\ \mathbf{x}\ $ (daily %)	ν	$P(z_t = z_{t-1})$
Normal	4,723	53.6%	0.98	6.2	0.994
Elevated	3,023	34.3%	2.03	3.9	0.991
Crisis	1,071	12.1%	2.09	5.5	0.993

3.5 Factor Pair Selection

We focus on the HML–SMB pair for three reasons. First, *economic plausibility*: Value and Size factors have documented overlap in institutional holdings (Section 4.4), providing an economically motivated mechanism—when leveraged institutions must unwind value positions, the associated size exposure is simultaneously deleveraged, creating a predictive lag from HML to SMB. This economic hypothesis precedes and motivates our data analysis. Second, *in-sample screening*: a preliminary screen of all 30 directed factor pairs found that HML–SMB shows the most distinctive regime-heterogeneous pattern, consistent with the economic prior. Crucially, the primary (LL-only) fit independently identifies HML–SMB as distinctive (Normal: $p = 1.0 \times 10^{-6}$, Elevated: $p = 0.013$, Crisis: $p = 0.880$ at lag 5)—no 2008-based screening is involved in pair selection. Third, *architectural contrast*: lag-1 evidence is strongest in Normal and substantially weaker outside Normal under the Bonferroni-significant threshold.

Pair selection and multiple comparisons. The OOS frozen test (Section 5.2) is treated as an *economically motivated* corroborative test of the HML–SMB pair (identified through both economic prior and in-sample screening; formal pre-registration was not performed), not a blind search among 30 pairs. The economic motivation (liquidity-mediated HML→SMB contagion) was established prior to the OOS evaluation and is consistent with the FF25 mechanism evidence (Section 4.5). Readers should nonetheless note that applying a 30-pair Bonferroni penalty to the OOS HAC $p = 0.041$ would yield a corrected $p > 1$; the OOS result’s strength rests on the economic motivation for the HML–SMB pair, not on the unadjusted p -value alone. Additionally, adjusting only for the 3 regimes (not the 30 pairs) gives a threshold of $\alpha/3 = 0.0167$; the F - $p = 0.014$ result barely survives this regime-adjusted threshold while the HAC- $p = 0.041$ does not.

4 Results

4.1 Regime Characteristics

Table 1 summarizes the three identified regimes. Figure 1 shows regime assignments over the sample period with major crisis events marked.




Figure 1: Regime assignments over sample period (1990–2024). Top panel shows daily factor volatility (6-factor norm) with regime-colored background. Bottom panel shows regime classifications. Vertical dashed lines mark major stress events. Under the primary (leakage-safe) fit, Crisis regime (red) reflects a high-kurtosis statistical state; the 2008-screened sensitivity fit aligns it with calendar crises (see Section 4.2). Elevated regime (yellow) captures moderate volatility periods.

Table 2: Crisis Detection Comparison. Detection rate = proportion of event days assigned to Crisis regime.

Event	Period	Student- t	Gaussian
2008 Financial	Jul '08 – Jun '09	0.0%	83.3%
2011 EU Debt	Jul – Oct 2011	0.0%	25.9%
2020 COVID-19	Feb – Jun 2020	0.0%	81.7%

4.2 Gaussian vs. Student- t Comparison

Under the leakage-safe primary selection rule, the Student- t fit does not label these benchmark windows as Crisis (0.0%), while the Gaussian baseline assigns substantial Crisis mass to 2008 and 2020. We therefore treat event-level “detection” as sensitivity-dependent and avoid confirmatory claims based solely on this table. **Terminology note.** Under the leakage-safe primary fit, the third regime captures a distinct statistical state (high kurtosis, shifted covariance) rather than exclusively named crisis events—Table 2 shows 0% alignment with 2008, 2011, and 2020 stress windows. We retain the label “Crisis” for consistency with the HMM literature [10], but readers should interpret Crisis-regime results as pertaining to this statistical state. The 2008-screened sensitivity fit does align with calendar crises; we report it separately where noted.

4.3 Main Result: Regime-Dependent Predictive Structure

Key finding under Bonferroni-significant threshold: only Normal-regime HML→SMB survives Bonferroni correction.

- **Normal (Bonferroni-significant):** Strong lag-1 effect ($p = 8.75 \times 10^{-9}$), robust to HAC correction. The Granger test conditions on lagged SMB, so the HML coefficient reflects *incremental cross-factor* predictive content beyond SMB’s own history—not a within-SMB autocorrelation artifact.

Table 3: Granger Causality Between HML and SMB by Regime. Significance threshold: $p < 0.00033$ (Bonferroni-corrected for 30 factor pairs).

Regime	Direction	p -value	Lag	Significant
Normal	HML $\rightarrow$ SMB	8.75×10^{-9}	1	Yes
Normal	SMB $\rightarrow$ HML	0.864	1	No
Elevated	HML $\rightarrow$ SMB	0.004	1	Suggestive*
Elevated	SMB $\rightarrow$ HML	0.897	1	No
Crisis	HML $\rightarrow$ SMB	0.695	1	No
Crisis	SMB $\rightarrow$ HML	0.294	1	No

*Below 0.01 but above Bonferroni threshold $\alpha/30 = 0.00033$.

Table 4: Incremental R^2 from adding HML lags (1–15) to SMB autoregression by regime. $\Delta R^2 = R^2_{\text{unrestricted}} - R^2_{\text{restricted}}$

Regime	n	R^2_{AR}	ΔR^2	p -value
Normal	4,361	0.61%	1.04%	5.98×10^{-5}
Elevated	2,690	0.59%	1.32%	0.002
Crisis	981	1.01%	0.56%	0.988

- **Elevated (suggestive):** HML $\rightarrow$ SMB is directional at conventional levels ($p = 0.004$) but does not pass the paper’s multiple-testing threshold.
- **Crisis (null):** No detectable lag-1 Granger effect in either direction.

Robustness to heteroskedasticity. Financial returns exhibit volatility clustering, which may inflate standard Granger F -test significance. We recompute at lag 1 (BIC-optimal) using Newey-West [46] heteroskedasticity-and-autocorrelation-consistent (HAC) standard errors. Normal-regime HML $\rightarrow$ SMB survives HAC correction ($p = 9.45 \times 10^{-8}$ HAC vs. $p = 8.75 \times 10^{-9}$ standard). Elevated HML $\rightarrow$ SMB weakens under HAC (full-sample in-sample: $p = 0.034$), and Crisis remains non-significant ($p = 0.711$ HAC vs. $p = 0.695$ standard).

Incremental R^2 . Table 4 quantifies the economic magnitude. Adding HML lags (1–15) to the SMB autoregression increases R^2 by 1.04% in Normal, 1.32% in Elevated, and 0.56% in Crisis. Normal and Elevated pass conventional significance levels ($p = 5.98 \times 10^{-5}$ and $p = 0.002$), while Crisis does not ($p = 0.988$). In the reverse direction, SMB lags add 1.35% in Crisis with $p = 0.562$, so we do not claim directional confirmation there.

All-pairs perspective. Figure 2 shows Granger causality p -values for all directed factor pairs across three regimes (lag 5, for visual density across all 30 pairs; Table 3 uses BIC-selected lag 1). Several pairs show regime-dependent variation; MKT-RF $\rightarrow$ SMB has the largest differential. HML $\rightarrow$ SMB ranks fourth in this differential metric, with strongest evidence in Normal ($p = 1.0 \times 10^{-6}$ at lag 5), suggestive Elevated evidence ($p = 0.013$), and no Crisis evidence ($p = 0.880$). We therefore treat HML $\rightarrow$ SMB as one regime-sensitive relation among several, not the dominant pair.

Lag specification and persistent structure. We tested lags 1–15 for HML $\rightarrow$ SMB across regimes (Figure 6). In Normal, the effect is strongest at lag 1 and remains jointly significant when including




Figure 2: All-pairs Granger causality by regime ($-\log_{10}(p)$, lag=5). Bordered cells with asterisks survive Bonferroni correction. The strongest regime-differential pair is MKT-RF $\rightarrow$ SMB; HML $\rightarrow$ SMB appears as a secondary differential pattern.

up to lag 15 (joint F -test of lags 1–15: $p = 5.98 \times 10^{-5}$, Table 4). Elevated shows weak-to-moderate joint evidence (lags 1–15 joint $p = 0.002$), while Crisis is consistently null (lag 1 $p = 0.695$, joint lag 1–15 $p = 0.988$). We present Elevated/Crisis lag patterns as exploratory because they do not meet the Bonferroni-significant threshold.

4.4 Hypothesized Economic Mechanism

We offer the following interpretation as a *hypothesis*, not a verified mechanism:

Normal-regime lag-1 link (Bonferroni-significant): During calmer periods, cross-factor inventory management and short-horizon rebalancing may transmit value shocks into size returns with a one-day delay.

Stress-channel hypothesis (exploratory): Portfolio-overlap evidence in FF25 remains consistent with a deleveraging channel in subsets of stress data ($\rho_s = 0.35$, permutation $p = 0.046$), but primary Crisis-regime factor-level Granger tests are not significant in the leakage-safe fit.

Alternative explanations we cannot rule out:

- Differential price delay [36] between value and size constituents could generate mechanical daily cross-autocorrelation independent of any economic channel
- A common driver (e.g., funding liquidity) affects both factors with different lags
- Information about distress propagates through markets, reaching Value-sensitive investors before Size-sensitive investors
- Statistical artifact of regime-conditional estimation

Greenwood and Thesmar [30] document that stocks with concentrated institutional ownership exhibit correlated fragility during liquidation events—consistent with our hypothesized deleveraging channel. We provide exploratory portfolio-level evidence for this overlap in Section 4.5.

4.5 Portfolio Overlap Evidence (Sensitivity Fit, Seed 42)

Note on model fit: All results in this section use the sensitivity fit (seed 42), which assigns 47.8% of COVID-2020 to Crisis and is selected for its superior crisis-period detection; the primary Granger

results (Section 4.3) use seed 28. The two seeds correspond to distinct local optima with different regime assignments; results from this section do not transfer to the primary fit.

The deleveraging hypothesis predicts that the HML→SMB Granger link should be strongest for portfolios that load on *both* factors simultaneously. The Small-HighBM portfolio sits in both the SMB long leg (small stocks) and the HML long leg (high book-to-market), making it the primary candidate transmission channel.

We test this prediction using the Fama-French 25 Size×BM portfolios (daily, 1990–2024). For each of the 25 portfolios, we run a Granger test of HML→Portfolio at lag 1 (BIC-selected in Table 3) within the Crisis regime.⁴

Spatial gradient test (exploratory). We define an “overlap score” for each portfolio as (4 – size quintile) + BM quintile, ranging from 0 (Big/LowBM, no overlap) to 8 (Small/HighBM, maximal overlap). The pre-specified test is a Spearman rank correlation between overlap score and $-\log_{10}(p)$ of the Granger test across 25 portfolios, with significance assessed by spatial permutation (10,000 shuffles of grid assignments; time series are never touched).

Result: $\rho_s = 0.35$, permutation $p = 0.046$ (asymptotic Spearman $p = 0.086$; Kendall $\tau = 0.27$, $p = 0.065$). The permutation test accounts for spatial autocorrelation in the grid; asymptotic and Kendall tests are borderline non-significant, so this evidence should be treated as *suggestive rather than confirmatory* (Figure 3). Granger significance concentrates in small-cap portfolios: all five Small quintile portfolios are nominally significant ($p < 0.01$), with three surviving Bonferroni correction ($\alpha/25 = 0.002$, using family-wise error rate $\alpha = 0.05$ for this exploratory test): S/L ($p = 2.5 \times 10^{-6}$, $\Delta R^2 = 1.47\%$), S/3 ($p = 0.0014$), and S/H ($p \approx 0$, $\Delta R^2 = 1.79\%$). One Big-quintile portfolio is nominally significant ($p < 0.01$), but none survive Bonferroni. This spatial pattern is inconsistent with a uniform factor-level effect and supports the overlap channel.

Overlap fraction. At lag 15 (the longest lag in the specification, chosen to measure cumulative exposure after initial propagation), $\Delta R^2(\text{HML} \rightarrow \text{S/H}) / \Delta R^2(\text{HML} \rightarrow \text{SMB}) = 0.392$ (S/H $\Delta R^2 = 1.33\%$; HML→SMB $\Delta R^2 = 3.38\%$; both are single-lag-15 portfolio-level regressions within Crisis, distinct from the joint lags-1–15 factor-level estimate in Table 4): the Small-HighBM portfolio alone accounts for 39% of the HML→SMB lag-15 incremental R^2 .

4.6 Early Warning Assessment

The Student- t HMM detects crisis regimes before market peaks (Table 5). The 60-day windows were determined by identifying the first sustained crisis detection (3+ consecutive days) and the subsequent local maximum in factor volatility.

Under the primary fit, early warning is absent for EU Crisis 2011 and COVID 2020 (Table 5), consistent with the primary fit’s non-alignment to these calendar events (Table 2).

⁴This analysis uses the sensitivity fit (seed 42), which detects 47.8% of COVID-2020 as Crisis (vs. 0% under primary seed 28); seed 42 Crisis regime spans 1,641 days vs. 1,071 under seed 28. Note: in the 50-seed frozen OOS Granger analysis, seed 42 belongs to the lowest-LL local optimum (train LL≈137,712; Elevated F - $p = 0.466$, null). The FF25 portfolio-level test is run here as an exploratory mechanism check under the fit that maximizes Crisis-period detection; it does not carry over the Granger significance from the two higher-LL optima.




Figure 3: FF 25 Portfolio Overlap Analysis. Panel A: $-\log_{10}(p)$ for HML→Portfolio Granger tests (Crisis, lag=1). Significance concentrates in Small quintile; ** Bonferroni-significant. Panel B: Spatial gradient—overlap score vs. Granger significance ($\rho_s = 0.35$, permutation $p = 0.046$).

Table 5: Early Warning Lead Time. First Detection = first day of 3+ consecutive Crisis regime assignments.

Event	First Det.	Vol. Peak	Lead
Lehman 2008	May 9, 2008	Sep 15, 2008	129 days
EU Crisis 2011	None	Aug 8, 2011	None
COVID 2020	None	Mar 23, 2020	None

Table 6: Event-Based Consistency Checks (full-sample primary fit regime assignments). Windows: 2008 Financial Jul ’08–Jun ’09, 2011 EU Debt Jul–Oct 2011, 2015 China Aug–Oct 2015, 2018 Vol Shock Feb 2018, 2020 COVID Feb–Apr 2020 (83 Crisis-classified days), 2022 Rate Hikes Jan–Dec 2022. Expected pattern: HML → SMB significant during Crisis. Result codes: ✓ = expected ($p < 0.10$ for HML→SMB, $p > 0.10$ reverse); Dir. = correct direction but insufficient power; × = opposite.

Event	Days	$p(\text{H} \rightarrow \text{S})$	$p(\text{S} \rightarrow \text{H})$	Result
2008 Financial	189	0.118	0.648	Dir.
2011 EU Debt	85	0.090	0.321	✓
2015 China	64	0.419	0.459	Dir.
2018 Vol Shock	29	0.236	0.179	×
2020 COVID	83	0.062	0.134	✓
2022 Rate Hikes	188	0.626	0.112	×

5 Discussion

5.1 Event-Based Consistency Checks

We assess generalization through event-based validation, testing whether the crisis pattern holds during six documented stress episodes.

Summary (exploratory, low-power): The expected pattern (HML→SMB significant, reverse not) holds clearly in 2 of 6 events (2011, 2020). Note: three events (2015, 2020, 2022) also appear in the frozen OOS Table 8; differences in p -values reflect distinct HMM fits (full-sample vs. 1990–2012 frozen). Two additional events (2008, 2015) show the correct direction ($p_{\text{HML} \rightarrow \text{SMB}} < p_{\text{SMB} \rightarrow \text{HML}}$) but lack statistical power. Under the null hypothesis of no directional preference, the probability of 4+ correct-direction events out of 6 by chance is approximately 0.34 ($P(X \geq 4 \mid n=6, p=0.5) = 0.34$, binomial test), so the directional pattern alone provides limited statistical evidence against the null.

Table 7: Frozen OOS Granger Causality (HML→SMB), 2013–2024. HMM trained on 1990–2012 only; relabeling uses train-period ordering; lag = 1 pre-fixed from in-sample BIC; filtered probabilities (no future info). HAC- p uses Newey–West bandwidth 1 (matching Granger lag); Elevated HAC $p = 0.041$.

Regime	n	F - p	HAC- p	ΔR^2
Normal	661	0.149	0.153	0.32%
Elevated	836	0.014	0.041	0.73%
Crisis	1,299	0.281	0.290	0.09%

Exceptions. Two events show reversed dynamics: 2018 Vol Shock and 2022 Rate Hikes (SMB→HML direction). Both are short-duration or policy-driven episodes distinct from the liquidity-driven crises (2008, 2011, 2020) where the HML→SMB direction is supported. This distinction between crisis types is speculative—we lack direct evidence for the underlying mechanisms.

Elevated regime finding. In the primary fit (full-sample in-sample), Elevated shows nominal HML→SMB evidence at lag 1 ($p = 0.004$) but this does not survive the paper’s Bonferroni threshold ($\alpha/30 = 0.00033$), and reverse direction remains null ($p = 0.897$). Annual splits are mixed (67% of years show SMB→HML dominant, not tabulated), reinforcing that Elevated evidence is suggestive rather than confirmatory.

5.2 Frozen-Parameter Out-of-Sample Test (Primary Validation)

Two-stage approaches risk circularity: the HMM uses the same returns for regime assignment that Granger tests subsequently analyze. We resolve this by fitting the Student- t HMM exclusively on 1990–2012 ($n = 5,797$ days), freezing all parameters ($\mu_k, \Sigma_k, \nu_k, A$), and classifying 2013–2024 ($n = 3,020$ days) without refitting. Granger tests are then run on the held-out labels using the same `select_lag_bic + run_granger_at_lag` pipeline as Table 3, ensuring full methodological consistency.

The frozen HMM assigns 24.0% to Normal, 30.7% to Elevated, and 45.3% to Crisis in the held-out period (vs. 21.5%, 13.7%, 64.8% in training). Notably, the Elevated-regime share more than doubles from 13.7% (training) to 30.7% (test), consistent with the OOS Elevated significance. After lag-1 boundary exclusion (same rule as Section 3.4), the OOS retains 661+836+1,299 = 2,796 of 3,020 days (7.4% excluded), lower than the full-sample 22.2% because the frozen classifier produces longer regime runs in the OOS period (fewer transitions when parameters are fixed).

Result. HML→SMB shows suggestive evidence in the Elevated regime ($p = 0.014$; HAC $p = 0.041$; $\Delta R^2 = 0.73\%$ at lag 1), while Normal ($p = 0.149$) and Crisis ($p = 0.281$) are null. (Note: Table 4 reports $\Delta R^2 = 1.32\%$ for Elevated using the full-sample in-sample fit with lags 1–15; the OOS $\Delta R^2 = 0.73\%$ uses lag 1 only on the 2013–2024 held-out period.) This result does not survive Bonferroni correction across 30 pairs (threshold $\alpha/30 = 0.00033$, family-wise $\alpha = 0.01$); its evidential weight rests on the economic motivation for the HML–SMB pair (Section 3.5). Notably, in-sample significance is expected given the degrees of freedom consumed by regime identification, lag selection, and pair screening on the same data; the OOS

Table 8: Held-Out Stress Events (2013–2024). HMM trained on 1990–2012 only. Windows: 2015 China Aug–Oct 2015 (37 days), 2020 COVID Feb–Jun 2020 (94 Elevated-classified days; wider window than Table 6’s 83 Crisis days), 2022 Rate Hikes Jan–Dec 2022.

Event	Days	Crisis%	$p(H \rightarrow S)$	$p(S \rightarrow H)$
2015 China	37	54%	0.165	0.123
2020 COVID	94 (Elevated)	0%	0.032	0.107
2022 Rate Hikes	124	100%	0.472	0.186

result is therefore the more informative—if more modest—test of generalization. A full 50-seed sensitivity analysis finds that **2 of 3 local optima** (covering 41 of 50 seeds; F - p and HAC- $p < 0.05$; median $p = 0.0075$) show Elevated significance. The 50 seeds converge to three local optima (train LL≈138,002: 28 seeds, all significant; LL≈137,942: 13 seeds, all significant; LL≈137,712: 9 seeds, all null); significance arises in the two higher-LL optima, each representing a distinct HMM solution rather than an independent initialization. These results use forward-only (filtered) HMM probabilities and train-period relabeling, eliminating all forms of look-ahead and circularity. OOS is treated as an economically motivated corroborative test of the pre-identified pair (Section 3.5); the Bonferroni-corrected result is discussed therein. The reverse direction SMB→HML in Normal is nominally significant (F - $p = 0.047$; HAC- $p = 0.063$) but does not survive Bonferroni correction across 30 pairs and was not pre-specified; we report it for completeness only.

Regime interpretation. The in-sample Normal-regime dominance ($p = 8.75 \times 10^{-9}$) vs. OOS Elevated significance reflects regime redistribution under the frozen classifier: in training (1990–2012), Crisis accounts for 64.8% of days while Elevated is only 13.7%; in the held-out period (2013–2024), Crisis shrinks to 45.3% while Elevated grows to 30.7%. The OOS period thus contains substantially more Elevated-regime days than the training period, and the frozen classifier concentrates the HML→SMB signal in this regime. The directional finding—HML predicts SMB during stress-adjacent (Elevated) conditions—is consistent across the OOS design.

Event-level detail. Restricting to specific stress episodes within the held-out period:

Event-level behavior is heterogeneous (COVID $p = 0.032$; 2022 $p = 0.472$), highlighting that the per-regime pooled result captures the aggregate stress-period signal more reliably than individual events.

5.3 Markov-Switching VAR Comparison

Our two-stage procedure (fit HMM, then test Granger) invites comparison with Markov-Switching VARs [21, 40], which jointly estimate regime-dependent dynamics in a single stage.

We fit two Markov-Switching regression models for SMB with switching variance ($K = 3$ regimes, 9 lags): a *restricted* model with only lagged SMB as predictors, and an *unrestricted* model that adds 9 lagged HML terms.

Likelihood ratio test: The unrestricted model significantly improves fit (LR = 114.95, $df = 27$ [3 regimes $\times$ 9 HML lag coefficients; variance parameters identical across restricted/unrestricted], $p = 8.0 \times 10^{-13}$), confirming that HML lags contain information

Table 9: Trading Strategy Backtest (1995–2024). Strategy trades SMB based on 9-day lagged HML signal during Crisis regimes only.

Metric	Strategy	Buy-and-Hold SMB
Annual Return	−0.3%	+0.1%
Sharpe Ratio	−0.07	+0.06
Max Drawdown	−30.0%	−43.5%

about SMB dynamics not captured by SMB’s own history, even within a single-stage regime-switching framework.

BIC comparison: Despite the significant likelihood improvement, BIC favors the restricted model ($\Delta\text{BIC} = -130$), reflecting the penalty for 27 additional parameters across 3 regimes. This indicates that while HML lags are statistically informative, the improvement does not justify the model complexity under a parsimony criterion.

Implication. The MS-VAR confirms that the HML–SMB predictive link exists within a principled single-stage framework ($p < 10^{-12}$), but its economic magnitude is too small to justify a complex regime-switching prediction model—consistent with our $\Delta R^2 \approx 1\%$ finding and the trading strategy failure. The LR test is aggregate; the MS-VAR does not separately decompose per-regime HML significance, so the two-stage regime-conditional Granger results (Table 3) remain the primary decomposition of where in the regime structure the predictive link concentrates.

5.4 Trading Strategy Evaluation

We evaluate whether the predictive relationship translates to profits via a simple strategy using the primary fit (seed 28): during Crisis regime, go long SMB when HML return over the past 9 days is positive, short SMB when negative. The Crisis-regime Granger test is null ($p = 0.695$), so we expect no edge; this test serves as a falsification check confirming that an absent statistical signal does not produce trading profits. Normal-regime predictability (lag 1) does not lend itself to a simple trading rule because it requires real-time regime classification and the lag-1 effect is small ($\Delta R^2 < 1\%$).

Why does Granger causality not imply trading profits?

The negative returns highlight an important distinction: *statistical predictability* $\neq$ *economic predictability*. Several factors explain the gap:

- (1) **Effect size:** The Granger test detects whether HML *improves* SMB prediction, not whether the improvement is large. The ΔR^2 from adding HML lags is 0.56% in the Crisis regime (Table 4); this is economically small and not statistically significant.
- (2) **Sign vs. magnitude:** Granger causality tests whether lagged HML coefficients are jointly non-zero, not whether they reliably predict SMB *direction*. The relationship may involve complex dynamics (e.g., mean reversion) that a simple directional strategy cannot capture.
- (3) **Regime detection lag:** The HMM assigns regimes using filtered probabilities (past data only). Real-time regime detection would be noisier, introducing additional error.
- (4) **Transaction costs:** Not modeled, but would further erode returns.

Table 10: Out-of-Sample VaR Backtest (2013–2024). Target violation rate: 5%. Christoffersen CC [16] tests conditional coverage (extending Kupiec’s [41] unconditional coverage test with a serial-independence component).

Model	Viol. Rate	Dev. from 5%	CC p -value
Unconditional	6.42%	+1.42 pp	0.002
GARCH(1,1)	3.91%	−1.09 pp	0.014
Regime-Conditional	6.69%	+1.69 pp	< 0.001
HML-Informed	5.77%	+0.77 pp	0.171
Hybrid (HMM+Vol)	5.60%	+0.60 pp	0.336

Hedging vs. alpha. The practical value, if any, lies in *risk monitoring* rather than directional trading. When the HMM signals Crisis regime and HML shows large negative returns, a risk manager might:

- Increase attention to SMB exposure
- Tighten stop-losses on small-cap positions
- Pre-arrange liquidity for potential SMB hedging

These are qualitative adjustments informed by the Elevated-regime predictive evidence (not the null Crisis Granger result, which serves as a falsification check), not mechanical trading rules. We formalize this intuition in Section 5.5.

5.5 Risk Monitoring Application (Sensitivity Fit, Seed 42)

Note on model fit: All VaR results in this section use the sensitivity fit (seed 42); the frozen OOS Granger result for seed 42 is null (Elevated $p = 0.466$), so VaR improvements reflect the volatility-override mechanism rather than direct Granger-signal transmission.

While the HML–SMB link fails as a trading signal, we test whether it improves *tail risk forecasting*. We compare five Value-at-Risk (VaR) models for SMB at the 5% level, trained on 1990–2012 and tested out-of-sample on 2013–2024 ($n \approx 3,000$ days; exact counts vary by model due to regime-boundary exclusion and warm-up requirements).⁵

- (1) **Unconditional:** 60-day rolling historical VaR.
- (2) **GARCH(1,1):** One-step volatility forecast with parametric VaR.
- (3) **Regime-conditional:** VaR window adapts to regime (Crisis: 30 days, Elevated: 50, Normal: 75).
- (4) **HML-informed:** Regime-conditional, plus a $2.3\times$ stress multiplier when 9-day cumulative HML < -0.5 during Crisis regime. Parameters calibrated on training data only.
- (5) **Hybrid (HMM+Vol):** HML-informed model plus a volatility override that forces Stress Alert when 20-day realized volatility exceeds the 95th percentile of the training period.

⁵We do not include CAViaR [23] (conditional autoregressive quantile regression) as a benchmark; our hybrid model combines regime detection with historical simulation rather than direct quantile autoregression, and the risk monitoring focus motivates simpler, interpretable specifications.

Both the HML-informed and hybrid models pass the Christoffersen conditional coverage test ($p > 0.05$); all other specifications are rejected. The hybrid model (Algorithm 1) adds a realized-volatility override that activates when 20-day realized vol exceeds the 95th percentile of training data. This corrects the frozen HMM’s COVID-2020 detection failure (0%→47.8% detection; Tables 11, 12) while improving VaR coverage to 5.60% ($p_{CC} = 0.336$). Note: these results use the sensitivity fit (seed 42), whose frozen OOS Granger result is null (Elevated $p = 0.466$); the VaR improvement therefore reflects the volatility-override mechanism rather than direct Granger-signal transmission.

Tick loss and model selection. Despite superior coverage, the HML-informed model has higher mean tick loss (1.05 vs. 0.90 baseline). This apparent contradiction arises because tick loss penalizes conservatism: a model with correct coverage but wider VaR incurs higher tick loss on the 95% of non-violation days. The Diebold-Mariano test, applied to tick loss, would rank the under-covering baseline above the correctly-covering HML-informed model—a known limitation of tick loss for comparing models at different coverage levels. We report both metrics for transparency: tick loss measures forecast sharpness, while Christoffersen CC measures correctness.

False alarms and model complexity. In the HML-informed model, the stress multiplier activates on 376 days (12.8% of the test period) with a 93.4% false alarm rate. In the hybrid model, alerts activate on 440 days (14.9%) with a 93.2% false alarm rate—i.e., on most alert days, no VaR violation occurs. This is partly a base-rate artifact: with a 5% violation target and 12.8% alert rate, at most 39% of alerts can be “true” (5/12.8), yielding a 61% false alarm floor. The 93.4% exceeds this floor but reflects the cost of correct tail coverage: alert-day VaR averages $2.51\times$ wider than non-alert VaR, achieving zero violations on alert days at the cost of unnecessary capital reserves on non-violation alerts. The alternative (GARCH) avoids regime labels entirely and achieves 3.91% violation rate with no false-alarm mechanism, but its over-coverage means capital is consistently over-reserved. In practice, the choice depends on whether the institution prefers fewer but louder alerts (HML-informed) or consistently conservative risk limits (GARCH). Both outperform unconditional VaR; neither dominates on all criteria. We note that the HML-informed model’s success may partly reflect conservative VaR widening during high-volatility regimes rather than the Granger link per se, since the OOS Granger effect is regime-dependent (Elevated $p = 0.014$; Normal $p = 0.149$; Crisis $p = 0.281$) rather than uniformly active across the held-out period.

COVID-19 common-mode failure and hybrid detector. The frozen HMM classifies the entire COVID period as non-Crisis (0% Crisis days), so the HML-informed stress multiplier never activates—a complete regime detection failure. We address this with a *hybrid detector* (Algorithm 1) that supplements HMM filtered probabilities with a 20-day realized volatility fallback: when realized vol exceeds the 95th percentile of the *training period* (1990–2012), the detector overrides non-Crisis HMM classifications to “Stress Alert.” The volatility threshold ($\sigma_{95} = 1.446$) is a fixed design choice locked before test evaluation—no search over thresholds on the holdout period.

Hybrid detector results. The volatility override activates on 44 of 92 COVID trading days (47.8% detection rate vs. 0% for HMM

Table 11: Hybrid Detector: COVID-19 Detection Sensitivity. Vol threshold = percentile of 20-day realized volatility computed on 1990–2012 training data. Primary result uses 95th percentile (bolded).

Percentile	Threshold	Override Days	Detection Rate
85th	0.924	85	92.4%
90th	1.098	76	82.6%
95th	1.446	44	47.8%
99th	2.122	25	27.2%

Table 12: Out-of-Sample VaR Backtest (2013–2024) with Hybrid Detector. Target violation rate: 5%. Rows show the unconditional baseline and the two HML-augmented specifications; GARCH(1,1) and Regime-Conditional results appear in Table 10.

Model	Viol. Rate	Dev.	CC p
Unconditional	6.42%	+1.42 pp	0.002
HML-Informed	5.77%	+0.77 pp	0.171
Hybrid (HMM+Vol)	5.60%	+0.60 pp	0.336




Figure 4: Hybrid detector performance. Panel A: realized volatility with 95th percentile threshold and override periods (COVID window highlighted). Panel B: SMB returns with unconditional, HML-informed, and hybrid VaR. Panel C: cumulative violations vs. 5% target.

alone; Table 11). VaR performance improves further: violation rate 5.60% (deviation from 5% target: +0.60 pp), Christoffersen CC $p = 0.336$ —the strongest conditional coverage pass among all specifications (Table 12; Figure 4). Average violation magnitude also improves (−0.275 vs. −0.305 for HML-Informed), indicating less severe tail breaches when they occur.

Algorithm 1 Real-Time Factor Risk Monitor

Require: Frozen HMM parameters $(\mu_k, \Sigma_k, v_k, A)$; vol threshold σ_{95} ; HML threshold τ ; stress multiplier m ; regime windows (w_0, w_1, w_2)

- 1: **for** each trading day t **do**
- 2: Compute filtered regime: $\hat{z}_t \leftarrow \arg \max_k P(z_t = k \mid \mathbf{x}_{1:t})$
- 3: Compute 20-day realized vol: $\hat{\sigma}_t \leftarrow \text{std}(\|\mathbf{x}\|_{t-19:t})$
- 4: **if** $\hat{z}_t \neq \text{Crisis}$ **and** $\hat{\sigma}_t > \sigma_{95}$ **then**
- 5: $\hat{z}_t \leftarrow \text{Crisis}$ (Stress Alert) {Volatility override; uses sensitivity fit (seed 42)}
- 6: **end if**
- 7: Set VaR window: $w \leftarrow w_{\hat{z}_t}$
- 8: Compute base VaR: $\text{VaR}_t \leftarrow Q_{0.05}(\text{SMB}_{t-w:t-1})$
- 9: **if** $\hat{z}_t = \text{Crisis}$ **and** $\sum_{\ell=1}^9 \text{HML}_{t-\ell} < \tau$ **then**
- 10: $\text{VaR}_t \leftarrow m \cdot \text{VaR}_t$ {HML stress multiplier}
- 11: **end if**
- 12: **return** $\hat{z}_t, \text{VaR}_t$
- 13: **end for**

Table 13: Algorithm 1 Parameter Values. All calibrated on training data (1990–2012) only. σ_{95} is the 95th percentile of the empirical distribution of 20-day realized volatility over 1990–2012 (mechanically defined). m and τ were selected by grid search ($m \in \{1.5, 2.0, 2.3, 2.5, 3.0\}$, $\tau \in \{-0.3, -0.5, -1.0\}$) to minimize training-period CC deviation from 5%. No test-period data was accessed during calibration.

Parameter	Symbol	Value
Vol threshold (95th pctile)	σ_{95}	1.446
HML cumulative threshold	τ	-0.5
Stress multiplier	m	2.3
Window (Normal)	w_0	75 days
Window (Elevated)	w_1	50 days
Window (Crisis)	w_2	30 days
HML lookback	—	9 days

5.6 Robustness

Temporal dynamics. Figure 5 presents rolling 3-year Granger p -values for HML→SMB, showing that the predictive link is not a static property, with peaks during calendar stress periods (2008–2009, 2020), though this unconditional rolling test does not distinguish regime-conditional effects. The signal is strongest during the 2008–2009 and 2020 crises and near-absent during calm periods.

Lag specification. BIC selects lag 1 for all regime×direction combinations (Table 3). Figure 6 shows significance is robust across lags 1–15 in Normal, while Crisis remains non-significant across lags.

Subsample stability. Splitting at 2008: Crisis HML→SMB pre-2008 ($n = 659$, $p = 0.335$) and post-2008 ($n = 322$, $p = 0.599$) are both non-significant, so we avoid claims of subsample-confirmed Crisis effects.

Alternative regime methods. Threshold-based volatility regimes (realized volatility > 90th percentile = Crisis) also detect the pattern in sensitivity checks, but we treat this as exploratory given the primary-fit null in Crisis at lag 1.




Figure 5: Rolling 3-year unconditional Granger causality ($-\log_{10}(p)$, HML→SMB, lag=9). Background shading shows HMM regime assignments for reference. Predictive strength varies episodically, with peaks during periods overlapping known market stress events. This rolling analysis does not condition on regime labels; regime-conditional results are in Table 3.




Figure 6: Granger causality p -values for HML→SMB across lags 1–15 by regime. Normal regime (green) shows concentrated short-lag significance ($p < 10^{-7}$ at lag 1). Crisis regime (red) remains non-significant across the tested lag range.

Practitioner baselines (exploratory). Standard unconditional tools detect the HML→SMB asymmetry but miss the architectural distinction. Rolling 250-day Granger tests find HML→SMB significant in 36.2% of windows vs. 9.6% for the reverse—with concentration in Normal (47.1% vs. 13.7% reverse), lower rates in Elevated (21.3%), and mixed Crisis windows (27.7%). FEVD spillover is regime-insensitive in direction, failing to detect the dormant Elevated regime or the Normal/Crisis directionality shift. Regime conditioning does not merely sharpen detection—it reveals qualitatively different predictive architectures that rolling methods average away.

Weekly aggregation. Using weekly returns, the HML→SMB crisis pattern remains non-significant ($p = 0.114$, $n = 221$), consistent with weak evidence outside Normal.

Regime transitions (exploratory). Analyzing 60-day windows around Normal→Crisis transitions, the HML→SMB relationship strengthens upon crisis entry (median p : 0.058 → 0.021), suggesting discrete intensification of the predictive link.




Figure 7: Four-model diagnostic protocol: MSE improvement (%) from including HML lags, by regime and model class. Only linear models (blue) achieve significant improvement in Normal and Elevated; all nonlinear methods (RF, MLP, LSTM) fail to improve prediction.

Table 14: Four-Model Diagnostic Protocol: HML→SMB by Regime. Linear column shows MSE improvement (%); others show permutation p -values.

Regime	n	Linear	RF p	MLP p	LSTM p
Normal	4,496	0.86%**	0.69	0.20	0.63 [†]
Elevated	2,792	0.92%**	1.00	1.00	0.93 [†]
Crisis	1,017	0.43%**	0.13	0.96	0.83 [†]

** $p < 0.01$ (not Bonferroni-corrected). [†]LSTM: sensitivity fit (seed 42), 100 perms.

Model complexity diagnostic protocol. We apply four model classes—Linear (OLS), RF (100 trees), MLP (64-32), LSTM (32 hidden, 9-step)—to characterize the *complexity boundary* of the predictive relationship. Each predicts SMB from its own lags (restricted) vs. both SMB and HML lags (unrestricted), with permutation testing (200 shuffles for Linear/RF/MLP, 100 for LSTM; at 200 permutations, p -value SE ≈ 0.007 at $p = 0.01$) and expanding-window CV.

Table 14 and Figure 7 present the results. Linear improvements are modest (Normal 0.86%, Elevated 0.92%, Crisis 0.43%). However, *no nonlinear method improves prediction in any regime*: RF, MLP, and LSTM permutation tests are uniformly non-significant.⁶ This complexity characterization supports a conservative conclusion: incremental predictive content in the HML→SMB direction is captured by linear terms, though TE reveals residual nonlinear structure in the reverse direction (Table 15).

Transfer entropy triangulation. As a model-free check, we compute transfer entropy [28, 50] (Frenzel–Pompe kNN, $k = 5$, 200 permutations). Normal shows significant bidirectional TE (HML→SMB $z = 2.45$, $p = 0.007$; SMB→HML $z = 5.37$, $p < 0.005$ empirical bound with 200 permutations; note the minimum achievable empirical p is $1/201 \approx 0.005$, and the z -score is computed under a Gaussian approximation that may not hold in heavy tails).

⁶The LSTM was estimated under a separate HMM initialization (see footnote 1); qualitative conclusions are consistent across fits.

Table 15: Transfer Entropy: Multi-Lag (1–9) by Regime. z -scores are standardized against the permutation mean/SD (200 shuffles); p -values use the standard normal CDF approximation. Note: with 200 permutations the minimum empirical p is 0.005; $p < 10^{-6}$ reflects the normal-approximation tail.

Regime	n	HML→SMB z	p	SMB→HML z	p
Normal	4,496	2.45**	0.007	5.37***	$< 0.005^{\dagger}$
Elevated	2,792	2.41**	0.008	1.65*	0.049
Crisis	1,017	1.01	0.157	1.22	0.111

*** $p < 0.005$ ([†]empirical bound; 200 perms), ** $p < 0.01$, * $p < 0.05$.
 n = clean observations after boundary exclusion.

Elevated is weakly bidirectional at conventional levels (HML→SMB $p = 0.008$; SMB→HML $p = 0.049$), while Crisis is non-significant in both directions. The reverse-direction Normal TE signal appears despite non-significant reverse linear Granger ($p = 0.864$, Table 3). This bidirectional pattern is consistent with information flowing in both directions through different channels: Granger causality tests *incremental* predictive content conditional on own lags (a conservative, linear criterion), while transfer entropy captures total information transfer including nonlinear dependence. The two measures are complementary, not contradictory—SMB→HML nonlinear information transfer can coexist with HML→SMB linear predictive precedence.

Filtered vs. smoothed regime probabilities. Smoothed probabilities use the full sample (forward and backward pass); filtered probabilities use only past data ($P(z_t = k \mid x_1, \dots, x_t)$), simulating real-time classification. Overall agreement is 95.9%. Under filtered regimes, the Crisis HML→SMB result remains non-significant ($p = 0.601$ vs. $p = 0.695$ smoothed), with filtered Crisis containing 1,079 days (vs. 1,071 smoothed). The finding is qualitatively consistent across classification methods.

5.7 Limitations

- (1) **Granger $\neq$ structural causality.** Our core finding is that HML *Granger-causes* SMB in the Normal regime—i.e., lagged HML improves SMB prediction conditional on lagged SMB. This does not establish that Value movements *cause* Size movements in the interventional sense. A common factor (liquidity, sentiment) could drive both with different lags. The “deleveraging cascade” interpretation is a plausible hypothesis, not a verified mechanism.
- (2) **Portfolio-level (not holdings-level) mechanism evidence.** Our FF 25 overlap analysis (Section 4.5) provides portfolio-level evidence for the deleveraging channel within the Crisis regime; the corresponding Normal-regime analysis (where factor-level Granger is Bonferroni-significant) is left for future work. We do not analyze 13F holdings data or fund flows for direct verification.
- (3) **Pair selection.** The OOS corroborative test focuses on the economically motivated HML→SMB pair (informed by both economic prior and in-sample screening of 30 pairs; formal pre-registration was not performed). While the economic hypothesis (liquidity-mediated contagion) provides a priori justification, the direction of the link was also observed in in-sample screening. The economic and data-based motivations

are not fully independent, and without pre-registration this caveat cannot be entirely eliminated.

- (4) **Regime redistribution and seed sensitivity.** The frozen-parameter OOS design (Section 5.2) resolves regime-identification circularity. However, the 1990–2012-trained classifier distributes OOS days differently than training: Crisis shrinks from 64.8% (train) to 45.3% (test), while Elevated grows from 13.7% to 30.7%, shifting the most significant OOS regime from in-sample Normal to OOS Elevated. A full 50-seed sensitivity analysis shows Elevated significance in 2 of 3 local optima (covering 41 of 50 seeds); the 9 null seeds share the lowest of three local optima, so robustness holds across distinct HMM solutions. The in-sample result ($p = 8.75 \times 10^{-9}$, all-seed robust) remains the higher-powered characterization. Additionally, regime *semantic drift* is a concern: a classifier trained on 1990–2012 (64.8% Crisis) and applied frozen to 2013–2024 (45.3% Crisis, 30.7% Elevated) may use the Elevated label to capture conditions that, in-sample, fell in Normal. The directional finding—HML predicts SMB in stress-adjacent conditions—is consistent across the in-sample (Normal) and OOS (Elevated) results, but these regime labels may not map to identical economic states.
- (5) **Regime boundary exclusion and sample size.** Boundary cleaning (1–15 lags excluded per table) reduces effective samples. Crisis contains 1,071 days in the primary fit, and Crisis-regime HML→SMB is null at lag 1 ($p = 0.695$).
- (6) **Subsample instability outside Normal.** Normal HML→SMB is robust, but Crisis remains non-significant in both pre-2008 and post-2008 subsamples ($p = 0.335$ and $p = 0.599$).
- (7) **Risk monitoring scope.** VaR improvement is modest (6.42%→5.60% with hybrid detector) and may partly reflect conservative widening during high-volatility regimes rather than the Granger link itself. The hybrid detector's volatility override addresses the regime classifier's COVID-19 failure but adds one additional parameter (95th percentile threshold).

6 Conclusion

The leakage-safe primary fit supports the following conclusions. HML→SMB is strongly significant in the Normal regime at lag 1 ($p = 8.75 \times 10^{-9}$; HAC $p = 9.45 \times 10^{-8}$). Elevated evidence is suggestive but below the Bonferroni-significant threshold, and Crisis evidence is null at lag 1 ($p = 0.695$). The frozen out-of-sample Elevated-regime test is suggestive at conventional levels ($p = 0.041$, HAC; not Bonferroni-significant across 30 pairs), and event-level results remain mixed.

Does establish:

- Bonferroni-significant Normal-regime HML→SMB predictive precedence
- Exploratory FF25 overlap evidence ($\rho_s = 0.35$, permutation $p = 0.046$; sensitivity fit, seed 42)
- No evidence of tradable alpha (Sharpe = −0.07)
- Practical utility of a hybrid detector for VaR calibration (5.60% violations, CC $p = 0.336$)

Does not establish:

- Uniform significance across every regime
- Uniform event-level support across all stress episodes
- Structural causality (common drivers could explain the pattern)
- Trading profitability

Practical value: A hybrid VaR model combining HMM regime detection with a realized-volatility fallback achieves the closest coverage to the 5% target (5.60%, $p_{CC} = 0.336$), detecting 47.8% of COVID-2020 where the HMM alone detected 0%. Algorithm 1 formalizes the monitoring protocol with all parameters calibrated on training data.

For practitioners, the key insight is that factor lead-lag relationships are not merely stronger or weaker across regimes. In the recomputed primary fit, the clearest actionable signal is Normal-regime lag-1 HML→SMB. Crisis-regime monitoring is better treated as stress-sensitivity analysis than as a statistically established directional signal. Real-time implementation still benefits from the volatility override: the frozen HMM alone classified 0% of COVID-2020 as Crisis, while the hybrid detector reached 47.8% detection with one pre-specified threshold parameter.

Funding

This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

References

- [1] Donald WK Andrews. 1991. Heteroskedasticity and Autocorrelation Consistent Covariance Matrix Estimation. *Econometrica* 59, 3 (1991), 817–858.
- [2] Andrew Ang and Geert Bekaert. 2002. International Asset Allocation With Regime Shifts. *Review of Financial Studies* 15, 4 (2002), 1137–1187.
- [3] Andrew Ang and Joseph Chen. 2002. Asymmetric correlations of equity portfolios. *Journal of Financial Economics* 63, 3 (2002), 443–494.
- [4] Miguel Anton and Christopher Polk. 2014. Connected stocks. *Journal of Finance* 69, 3 (2014), 1099–1127.
- [5] Clifford S Asness, Tobias J Moskowitz, and Lasse Heje Pedersen. 2013. Value and momentum everywhere. *Journal of Finance* 68, 3 (2013), 929–985.
- [6] Pedro Barroso and Pedro Santa-Clara. 2015. Momentum Has Its Moments. *Journal of Financial Economics* 116, 1 (2015), 111–120.
- [7] Monica Billio, Mila Getmansky, Andrew W Lo, and Liorana Pelizzon. 2012. Econometric measures of connectedness and systemic risk in the finance and insurance sectors. *Journal of Financial Economics* 104, 3 (2012), 535–559.
- [8] Markus K Brunnermeier. 2009. Deciphering the liquidity and credit crunch 2007–2008. *Journal of Economic Perspectives* 23, 1 (2009), 77–100.
- [9] Markus K Brunnermeier and Lasse Heje Pedersen. 2009. Market Liquidity and Funding Liquidity. *Review of Financial Studies* 22, 6 (2009), 2201–2238.
- [10] Jan Bulla. 2011. Hidden Markov models with t components. Increased persistence and other aspects. *Quantitative Finance* 11, 3 (2011), 459–475.
- [11] Mark M Carhart. 1997. On Persistence in Mutual Fund Performance. *Journal of Finance* 52, 1 (1997), 57–82.
- [12] Maddalena Cavicchioli. 2025. Forecasting Markov Switching Vector Autoregressions: Evidence from Simulation and Application. *Journal of Forecasting* 44, 1 (2025), 136–152. doi:10.1002/for.3180
- [13] Ludwig B. Chincarini, Renato Lazo-Paz, and Fabio Moneta. 2026. Crowded Spaces and Anomalies. *Journal of Banking & Finance* 182 (2026), 107579. doi:10.1016/j.jbankfin.2025.107579
- [14] Francesco Chincoli and Massimo Guidolin. 2017. Linear and Nonlinear Predictability in Investment Style Factors: Multivariate Evidence. *Journal of Asset Management* 18, 6 (2017), 476–509. doi:10.1057/s41260-017-0048-5
- [15] Tarun Chordia and Lakshmanan Shivakumar. 2002. Momentum, Business Cycles, and Time-Varying Expected Returns. *Journal of Finance* 57, 2 (2002), 985–1019.
- [16] Peter F Christoffersen. 1998. Evaluating Interval Forecasts. *International Economic Review* 39, 4 (1998), 841–862.
- [17] Rama Cont. 2001. Empirical properties of asset returns: stylized facts and statistical issues. *Quantitative Finance* 1, 2 (2001), 223–236.
- [18] Joshua Coval and Erik Stafford. 2007. Asset Fire Sales (and Purchases) in Equity Markets. *Journal of Financial Economics* 86, 2 (2007), 479–512.

- [19] Kent Daniel and Tobias J Moskowitz. 2016. Momentum Crashes. *Journal of Financial Economics* 122, 2 (2016), 221–247.
- [20] Francis X Diebold and Kamil Yilmaz. 2012. Better to give than to receive: Predictive directional measurement of volatility spillovers. *International Journal of Forecasting* 28, 1 (2012), 57–66.
- [21] Matthieu Droumaguet, Anders Warne, and Tomasz Woźniak. 2017. Granger Causality and Regime Inference in Markov Switching VAR Models with Bayesian Methods. *Journal of Applied Econometrics* 32, 4 (2017), 802–818.
- [22] Sina Ehsani and Juhani T Linnainmaa. 2022. Factor Momentum and the Momentum Factor. *Journal of Finance* 77, 3 (2022), 1877–1919. doi:10.1111/jofi.13131
- [23] Robert F. Engle and Simone Manganelli. 2004. CAViaR: Conditional Autoregressive Value at Risk by Regression Quantiles. *Journal of Business & Economic Statistics* 22, 4 (2004), 367–381.
- [24] Eugene F Fama and Kenneth R French. 1993. Common risk factors in the returns on stocks and bonds. *Journal of Financial Economics* 33, 1 (1993), 3–56.
- [25] Eugene F Fama and Kenneth R French. 2015. A five-factor asset pricing model. *Journal of Financial Economics* 116, 1 (2015), 1–22.
- [26] Volker Flögel, Christian Schlag, and Claudia Zunft. 2022. Momentum-Managed Equity Factors. *Journal of Banking & Finance* 137 (2022), 106251.
- [27] Kenneth R French. 2024. Kenneth French Data Library. https://mba.tuck.dartmouth.edu/pages/faculty/ken.french/data_library.html. Accessed: 2024–2025.
- [28] Stefan Frenzel and Bernd Pompe. 2007. Partial Mutual Information for Coupling Analysis of Multivariate Time Series. *Physical Review Letters* 99, 20 (2007), 204101.
- [29] Clive WJ Granger. 1969. Investigating causal relations by econometric models and cross-spectral methods. *Econometrica* 37, 3 (1969), 424–438.
- [30] Robin Greenwood and David Thesmar. 2011. Stock price fragility. *Journal of Financial Economics* 102, 3 (2011), 471–490.
- [31] Massimo Guidolin and Allan Timmermann. 2007. Asset allocation under multivariate regime switching. *Journal of Economic Dynamics and Control* 31, 11 (2007), 3503–3544.
- [32] Tarun Gupta and Bryan Kelly. 2019. Factor Momentum Everywhere. *Journal of Portfolio Management* 45, 3 (2019), 13–36.
- [33] Valentin Haddad, Serhiy Kozak, and Shrihari Santosh. 2020. Factor Timing. *Review of Financial Studies* 33, 5 (2020), 1980–2018.
- [34] James D Hamilton. 1989. A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica* 57, 2 (1989), 357–384.
- [35] Alain Heq, Luca Margaritella, and Stephan Smeekes. 2023. Granger Causality Testing in High-Dimensional VARs: A Post-Double-Selection Procedure. *Journal of Financial Econometrics* 21, 3 (2023), 915–958. doi:10.1093/jfinc/nbab023
- [36] Kewei Hou and Tobias J. Moskowitz. 2005. Market Frictions, Price Delay, and the Cross-Section of Expected Returns. *Review of Financial Studies* 18, 3 (2005), 981–1020.
- [37] Clint Howard, Harald Lohre, and Sebastiaan Mudde. 2025. Causal Network Representations in Factor Investing. *Intelligent Systems in Accounting, Finance and Management* 32, 1 (2025), e70001.
- [38] Narasimhan Jegadeesh and Sheridan Titman. 1993. Returns to Buying Winners and Selling Losers: Implications for Stock Market Efficiency. *Journal of Finance* 48, 1 (1993), 65–91.
- [39] Amir E Khandani and Andrew W Lo. 2011. What happened to the quants in August 2007? Evidence from factors and transactions data. *Journal of Financial Markets* 14, 1 (2011), 1–46.
- [40] Hans-Martin Krolzig. 1997. *Markov-Switching Vector Autoregressions: Modelling, Statistical Inference, and Application to Business Cycle Analysis*. Springer, Berlin, Germany.
- [41] Paul H. Kupiec. 1995. Techniques for verifying the accuracy of risk measurement models. *Journal of Derivatives* 3, 2 (1995), 73–84.
- [42] Andrew W Lo and A Craig MacKinlay. 1990. When Are Contrarian Profits Due to Stock Market Overreaction? *Review of Financial Studies* 3, 2 (1990), 175–205.
- [43] Marcos M. López de Prado. 2023. *Causal Factor Investing: Can Factor Investing Become Scientific?* Cambridge University Press, Cambridge, UK. doi:10.1017/9781009397315
- [44] Dong Lou and Christopher Polk. 2022. Comomentum: Inferring arbitrage activity from return correlations. *Review of Financial Studies* 35, 7 (2022), 3272–3302.
- [45] Joseph M Marks and Chenguang Shang. 2019. Factor Crowding and Liquidity Exhaustion. *Journal of Financial Research* 42, 1 (2019), 147–180.
- [46] Whitney K Newey and Kenneth D West. 1987. A Simple, Positive Semi-Definite, Heteroskedasticity and Autocorrelation Consistent Covariance Matrix. *Econometrica* 55, 3 (1987), 703–708.
- [47] Lennart Oelschläger, Timo Adam, and Rouven Michels. 2024. fhMM: Hidden Markov Models for Financial Time Series in R. *Journal of Statistical Software* 109, 9 (2024), 1–35. doi:10.18637/jss.v109.i09
- [48] Judea Pearl. 2009. *Causality: Models, Reasoning, and Inference* (2nd ed.). Cambridge University Press, Cambridge, UK.
- [49] Zacharias Psaradakis, Morten O Ravn, and Martin Sola. 2005. Markov Switching Causality and the Money-Output Relationship. *Journal of Applied Econometrics* 20, 5 (2005), 665–683.
- [50] Thomas Schreiber. 2000. Measuring information transfer. *Physical Review Letters* 85, 2 (2000), 461.
- [51] Andrei Shleifer and Robert W Vishny. 1997. The Limits of Arbitrage. *Journal of Finance* 52, 1 (1997), 35–55.
- [52] Ali Shojaie and Emily B. Fox. 2022. Granger Causality: A Review and Recent Advances. *Annual Review of Statistics and Its Application* 9 (2022), 289–319. doi:10.1146/annurev-statistics-040120-010930
- [53] Yizhan Shu and John M. Mulvey. 2025. Dynamic Factor Allocation Leveraging Regime-Switching Signals. *Journal of Portfolio Management* 51, 3 (2025). doi:10.3905/jpm.2024.1.649
- [54] Yizhan Shu, Chenyu Yu, and John M. Mulvey. 2024. Downside Risk Reduction Using Regime-Switching Signals: A Statistical Jump Model Approach. *Journal of Asset Management* 25, 5 (2024), 493–507.
- [55] Jeremy C Stein. 2009. Presidential address: Sophisticated investors and market efficiency. *Journal of Finance* 64, 4 (2009), 1517–1548.
- [56] Alex Tank, Ian Covert, Nicholas Foti, Ali Shojaie, and Emily B Fox. 2022. Neural Granger Causality. *IEEE Transactions on Pattern Analysis and Machine Intelligence* 44, 8 (2022), 4267–4279.
- [57] Matthew Wang, Yi-Hong Lin, and Ilya Mikhelson. 2020. Regime-Switching Factor Investing with Hidden Markov Models. *Journal of Risk and Financial Management* 13, 12 (2020), 311.
- [58] Ivo Welch and Amit Goyal. 2008. A Comprehensive Look at the Empirical Performance of Equity Premium Prediction. *Review of Financial Studies* 21, 4 (2008), 1455–1508.